



1. FastAPI (for Agentic Workspace)

You must know:

- How FastAPI handles requests internally (ASGI, Starlette)
- Dependency injection
- Async vs sync endpoints
- Handling file uploads & streaming (SSE — which you used!)
- Background tasks
- How middlewares work
- Pydantic models + validation
- Rate limiting patterns (even high-level)

Interview-level questions:

- How does FastAPI achieve high performance?
- When to use async vs sync DB calls?
- How would you scale a FastAPI app?
- How does SSE work in FastAPI?

Since your RAG system uses streaming, they WILL ask about SSE.



2. PostgreSQL + pgvector

This is **very important** because pgvector + hybrid search is rare for freshers; it gives you a huge advantage.

You must know:

- How vectors are stored in Postgres
- What an index (IVFFlat/HNSW) is
- How approximate vector search works
- Why hybrid search (BM25 + vector) is needed
- How you did cross-encoder re-ranking

Interview-level questions:

- Why use pgvector instead of ElasticSearch?
 - Difference between cosine similarity and dot product?
 - How is a vector index built?
 - How do you prevent slow vector queries?
-



3. Redis (for URL Shortener)

Redis is *the most important* technology for backend interviews.

You must know:

- In-memory storage architecture
- Data types (string, hash, list, set, sorted set)
- Expiry mechanism
- Eviction policies
- Redis persistence (AOF, RDB)
- Redis atomic counters → **You used this**
- Pipelining & transactions
- Caching strategies (write-through, write-back, write-around)

Interview-level questions:

- How does Redis guarantee atomicity for counters?
- What happens if Redis crashes?
- How do you prevent cache stampede?
- Why Redis over DB indexing?

Since your project explicitly says “85% load reduction” — expect deep questions.



4. Spring Boot (for LMS)

This is another high-value topic.

You must know:

- Bean lifecycle, dependency injection (IoC)
- REST controller flow
- Filters vs Interceptors
- Spring Security → especially JWT auth
- JPA/Hibernate basics
- Transaction isolation levels
- Pagination, batching
- How to design scalable APIs

Interview-level questions:

- How does @Transactional work internally?
 - How does JWT authentication work end-to-end?
 - What is the difference between filter and interceptor?
 - How to scale a Spring Boot app horizontally?
-



5. Docker + AWS (you used EC2, S3, RDS)

You don't need deep DevOps knowledge, but you **must** know:

For Docker:

- Dockerfile basics
- Container lifecycle
- Image layers
- Why containerization makes deployment easier

For AWS:

- EC2 → Deploying backend
- RDS → Performance + scaling basics
- S3 → File storage + presigned URLs
- Security groups
- Load balancers (high-level)

Interview questions:

- What is the difference between VM and container?
 - How would you deploy your LMS system on AWS?
 - Why use S3 instead of storing files in database?
-



How Deep Should You Study Each Tech?

Not too deep — no need to know kernel-level internals.

You need:

Enough depth to defend design choices + answer “why” questions.

For example:

Skill	Level needed
FastAPI	Interview-ready + architecture awareness
pgvector	Conceptual + indexing strategies
Redis	Medium-deep (highly asked in interviews)
Spring Boot	Medium-deep (company dependent)
JWT, Auth	Deep (critical in backend interviews)
AWS, Docker	Medium (basic deploy knowledge)



The Most Important Areas (Ranked for Interviews)

If you have limited time, focus like this:

- ① Caching + Redis internals
- ② Authentication + JWT + RBAC
- ③ System design for your 3 projects
- ④ Postgres + indexing (incl. pgvector)
- ⑤ FastAPI internals (async, DI, SSE)
- ⑥ Spring Boot internals
- ⑦ AWS + Docker basics

If you master the above → you will 100% crack backend roles.

Must Know AI/ML interview questions -

1. RAG (Retrieval-Augmented Generation) — GUARANTEED Questions

Since your project uses hybrid search + pgvector + agents, expect heavy RAG questions.

Mandatory concepts you must be able to explain:

- ♦ What is RAG? Why use it instead of directly prompting the LLM?

Explain: hallucination reduction, grounding, domain adaptation.

- ♦ Retrieval pipeline

Be able to describe:

Document ingestion → Chunking → Embedding → Indexing → Retrieval → Re-ranking → Prompt construction → LLM response

- ◆ **Chunking strategies**

They may ask:

- What chunk size did you use?
- Why overlapping?
- How chunking affects retrieval quality?

- ◆ **Embeddings**

- What embedding model did you use? Why?
 - Cosine similarity vs dot product
 - Dimensionality trade-offs
 - How embedding quality affects retrieval
-



2. Vector Databases + pgvector — HIGH PROBABILITY

Because your backend uses **PostgreSQL + pgvector**, interviewers will dig here.

They may ask:

- How are vectors stored in pgvector?
 - What is IVFFlat or HNSW?
 - How approximate nearest neighbor (ANN) search works?
 - How to choose nlist, nprobe?
 - How do indexes speed up vector search?
 - When to use vector search vs keyword search (BM25)?
 - Why hybrid search is needed?
-



3. Hybrid Search (BM25 + Vector Search) — VERY HIGH PROBABILITY

Your resume explicitly mentions:

“hybrid search (vector + BM25) with cross-encoder re-ranking.”

Expect questions like:

◆ **Why use hybrid search?**

Explain:

- BM25 handles exact lexical matches
- Vector search handles semantics
- Combining improves both relevance + recall

◆ **How do you combine scores from BM25 + vector similarity?**

Talk about:

- Weighted sum
- Reciprocal rank fusion
- Normalization

◆ **What is cross-encoder re-ranking?**

Explain:

- Bi-encoder → fast, approximate
 - Cross-encoder → slow but highly accurate (token-level attention)
 - Used only on the top-k retrieved results
-



4. LangChain Agents — MEDIUM PROBABILITY

You implemented:

- AI agents for email drafting, Jira ticket creation, report generation

They may ask:

- What is an agent in the LangChain framework?
- How is an agent different from a chain?

- What tools did your agent use?
- How do you prevent infinite loops?
- How do you ensure safe tool invocation?

Bonus:

- ReAct prompting
 - Function calling vs tools
 - Agent memory
-



5. LLM Internals & Prompt Engineering — MEDIUM PROBABILITY

Not very deep — just conceptual.

You must know:

- What is a transformer? (High-level)
- What is attention? (Explain simply)
- Differences between:
 - **Prompting**
 - **Fine-tuning**
 - **RAG**
 - **Adapters (LoRA)**

Expect questions like:

- Why did you choose RAG over fine-tuning?
 - How do you reduce hallucinations?
 - How do you design prompts for your system?
-



6. Evaluation of RAG Systems — INTERVIEWERS LOVE THIS

Most candidates have no idea about evaluation.

They may ask:

- How do you measure retrieval quality?
 - Recall@k
 - Precision
 - Hit rate
 - How do you test RAG correctness?
 - How do you evaluate hallucinations?
 - How do you benchmark latency?
-



7. Scalability & Architecture — HIGH PROBABILITY

Your RAG system includes:

- SSE streaming
- Multi-tenant design
- Hybrid search

Expect:

- How did you scale embedding generation?
 - How do you handle large documents?
 - How do you store metadata for chunks?
 - How does SSE work?
 - How would you scale pgvector?
 - How do you handle concurrent user queries?
-



8. Data Pipeline Questions

Since you built ingestion:

Questions they can ask:

- How do you parse PDFs?
 - How do you handle tables, images?
 - What happens if a huge file is uploaded?
 - How do you ensure async job execution?
 - Did you do offline indexing or on-the-fly?
-



9. Security + Multi-Tenancy (Rare but valuable)

Your resume says:

“multi-tenant RAG platform with JWT auth and role-based access.”

Expect:

- How did you separate embeddings per tenant?
 - How do you prevent data leakage?
 - How do you handle access control for documents?
 - Why JWT and not session-based auth?
-



10. Behavioral Questions Specific to AI Projects

Interviewers will ask:

“What was the most difficult challenge in building this AI system?”

You should talk about:

- Retrieval quality
- Latency
- Large file processing
- Accuracy tuning
- Prompt failure cases

“What is one thing you would improve in your RAG system?”

Say:

- Structured retrieval
- Add reranking
- Add caching for embeddings
- Move to dedicated vector DB (Weaviate, Pinecone)

“What failures did you face and how did you solve them?”

Talk about:

- Wrong retrieval
 - Multi-tenant isolation
 - Query latency
 - Long document chunking issues
-

SUMMARY: What AI/ML Topics You MUST Prepare (Your Checklist)

◆ **MUST KNOW (Top priority):**

- RAG pipeline end-to-end
- Chunking & embeddings
- pgvector indexing
- Hybrid search
- Cross-encoder re-ranking
- LangChain Agents
- Retrieval evaluation metrics
- Scaling/latency optimization

◆ **GOOD TO KNOW (Extra points):**

- Transformers high-level
- Attention mechanism
- Why RAG > fine-tuning
- Safety, hallucination reduction

◆ **DO NOT NEED:**

- ✗ Backpropagation
- ✗ CNNs, RNNs
- ✗ Deep learning math
- ✗ ML algorithms like SVM, random forest
- ✗ Statistical ML

Because your project is **LLM application architecture**, not ML research.